

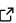
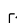
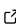
Processing raw mobility data with the checkin package

Michael Kane ^{1¶}, Owais Gilani ², and Simon Urbanek ³

¹ Biostatistics Department, School of Public Health, Yale University, New Haven, CT, USA ² Math Department, Bucknell University, Lewisburg, PA, USA ³ Statistics Department, University of Auckland, Auckland, NZ ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

The analysis of mobility data often begins with the processing of data capturing the time and location of individuals. When preparing these data for analysis, the continuous timestamp data can be particularly difficult to normalize since they require the data scientist to integrate mode assumptions such as, where is a person/device between checkins? Where is a person/device before the first checkin? Should last-observation-carried-forward be used between checkins or between discretization boundaries? For example, we may want to build sequences of checkins for a given device over locations and then count the total number of transitions between locations over all individuals. Problems like these are common in mobility research, they require careful consideration based on the goals of an analysis, and software tools implementing these types of computations will provide benefits in terms of time savings and data integrity. To address these challenges we provide the checkin package, which provides a standard set of functions for appropriately discretizing spatio-timestamp data for aggregate analysis for the R programming environment (R Core Team 2022).

Statement of Need

Raw mobility data is often characterized by having a column denoting the device/person identifier, a timestamp, the location, and potentially other features of the checkin. The device/person identifier is often given as a unique identifier of a device or person in possession of a device; the timestamp denotes the time at which a person was at a given location; and the location can be precise location information (such as GPS) but is often aggregated to a discrete location, such as a store, census tract, county etc. While these data are information-rich, to analyze them, especially at the aggregate level (many devices and many locations), requires processing to transform them into a representation amenable to analysis. One such representation is the *mobility graph*, which encodes vertices as discrete locations, directed edges as the aggregate movement between locations, and edge weights capturing the amount of movement (or similar measure) between locations with corresponding directed edge [[@gilani2020](#)]. More generally, processing steps for these aggregate analyses either discretize continuous timestamp and spatial data and/or aggregate already discretized data.

Basic operations for processing time-stamp are implemented in the core of R and there are a plethora of package that make operations such as reading, comparing, and adding offsets more convenient (see [Eddelbuettel \(2020\)](#), [Grolemund and Wickham \(2011\)](#), etc. for more) there has not been a set of standard functions specifically for processing checking data, which requires the following types of operations:

1. Find the location of person/device at a specified interval, using last-observation-carry forward if specified.
2. Construct a generator for creating sequences of intervals over which data should be processed.

43 3. A map operation working in conjunction with interval iterators for processing data over
44 spatio-interval data.

45 The checkin package provides an extensible library for providing these operations in a way that
46 is compatible with the “tidy data” approach described in Wickham (2014), it is compatible
47 with standard dplyr (Wickham et al. 2022) functions, and it uses foreach package (Microsoft
48 and Weston 2022) to provide a parallel-backend for compute-intensive computing. Data sets
49 returned from checkin functions are in the tibble (Grolemund and Wickham 2011) format.

50 Usage

51 Consider the checkins data set from the checkin library shown below. The data consists of 3
52 columns corresponding to the device (id), time, (timestamp), and location identifier (location).
53 There are a total of 1000 unique people/devices rangin in time from 2020-04-19 00:01:13 EST
54 to 2020-05-08 18:58:29 EST.

```
55 library(checkin)
56 library(dplyr)
57
58 ##
59 ## Attaching package: 'dplyr'
60
61 ## The following objects are masked from 'package:stats':
62 ##
63 ##   filter, lag
64
65 ## The following objects are masked from 'package:base':
66 ##
67 ##   intersect, setdiff, setequal, union
68
69 data(checkins)
70 checkins
71
72 ## # A tibble: 14,609 × 3
73 ##       id timestamp      location
74 ##   <int> <dtm>      <int>
75 ## 1  442 2020-04-19 17:23:44    36496
76 ## 2  442 2020-04-19 17:23:35    36496
77 ## 3  166 2020-04-19 12:55:44    37461
78 ## 4  476 2020-04-19 06:23:47    33476
79 ## 5  456 2020-04-19 05:39:09    33468
80 ## 6  458 2020-04-19 10:04:20    36500
81 ## 7  458 2020-04-19 15:15:25    36876
82 ## 8  651 2020-04-19 13:31:19    37391
83 ## 9  652 2020-04-19 05:27:54    37469
84 ## 10 653 2020-04-19 05:28:59    37389
85 ## # ... with 14,599 more rows
86 ## # Use `print(n = ...)` to see more rows
```

87 Now suppose we would like to examine the checkins of individual/device number 335. We will
88 extract all rows with id 335 and then specify an interval starting at the beginning of the hour
89 of the first checkin to one hour later using the checkins_in_interval() function. The code
90 and output are below and there are two things to note. First, the first row location is NA. This
91 is because the individual/device does not appear before 2020-04-19 00:41:46 and a location
92 cannot be determined. Second the original data set did not include an entry for id 335 at
93 2020-04-19 00:59:59. This value was carried forward from the previous location.

```

94 library(lubridate)
95
96 ##
97 ## Attaching package: 'lubridate'
98
99 ## The following objects are masked from 'package:base':
100 ##
101 ##    date, intersect, setdiff, union
102
103 x <- checkins %>%
104   filter(id == 335) %>%
105   arrange(timestamp)
106
107 start <- x$timestamp[1]
108 minute(start) <- 0
109 second(start) <- 0
110
111 end <- start + hours(1) - seconds(1)
112 checkins_in_interval(x, "timestamp", start, end)
113
114 ## # A tibble: 3 × 3
115 ##       id timestamp      location
116 ##   <int> <dtm>      <int>
117 ## 1    NA 2020-04-19 00:00:00      NA
118 ## 2   335 2020-04-19 00:41:46   32576
119 ## 3   335 2020-04-19 00:59:59   32576

```

Finally, suppose we would like to get the locations of individuals/devices at the beginning of a time interval, the location at the end of the interval, and the total amount of time the individual/devices was checked into the beginning location. This operation would be performed in two steps. First, a function, `from_to()` is constructed, which takes the rows corresponding to a single id for a given interval. This function finds, the initial location (`from`), the end location (`to`), the timestamp at the *beginning* of the interval, and the duration of the initial location. In the second step, we group the checkin data by id and pass the result to the `map_hourly_interval_dfr()` function, which applies to `from_to()` function to each id over each hourly interval. Other intervals are included in the package and the documentation includes a information on how to construct similar functions over custom intervals.

```

120
121
122
123
124
125
126
127
128
129
130 from_to <- function(it) {
131   it$duration <- c(diff(it$timestamp), 0)
132   units(it$duration) <- "secs"
133   it$duration <- as.numeric(it$duration)
134   from_duration <- sum(it$duration[it$location == it$location[1]])
135   tibble(from=it$location[1],
136          to = it$location[nrow(it)],
137          timestamp = it$timestamp[1],
138          from_duration = from_duration)
139 }
140
141
142 checkins |>
143   head(100) |>
144   group_by(id) |>
145   map_hourly_interval_dfr(from_to, time = "timestamp")
146
147 ## # A tibble: 119 × 5
148 ##       id      from      to timestamp      from_duration

```

```

148 ##      <chr> <int> <int> <dtm>                                <dbl>
149 ##    1 166    37461 37461 2020-04-19 13:00:00                    3599
150 ##    2 286    33165 33165 2020-04-19 18:00:00                    3599
151 ##    3 313    33149 33149 2020-04-19 19:00:00                    3599
152 ##    4 381    33359 33359 2020-04-19 16:00:00                    3599
153 ##    5 381    33359 33359 2020-04-19 17:00:00                    3599
154 ##    6 442    36496 36496 2020-04-19 18:00:00                    3599
155 ##    7 456    33468 33468 2020-04-19 06:00:00                    3599
156 ##    8 457    33130 33130 2020-04-19 13:00:00                    3599
157 ##    9 457    33130 33130 2020-04-19 14:00:00                    3599
158 ##   10 457    33130 33130 2020-04-19 15:00:00                    3599
159 ## # ... with 109 more rows
160 ## # Use `print(n = ...)` to see more rows

```

Acknowledgements

This work was partially supported by the National Science Foundation (NSF) Grant Human Networks and Data Science - Infrastructure (HNDS-I), award numbers 2024335 and 2024233.

References

- Eddelbuettel, Dirk. 2020. *Anytime: Anything to 'POSIXct' or 'Date' Converter*. <https://CRAN.R-project.org/package=anytime>.
- Grolemund, Garrett, and Hadley Wickham. 2011. "Dates and Times Made Easy with lubridate." *Journal of Statistical Software* 40 (3): 1–25. <https://www.jstatsoft.org/v40/i03/>.
- Microsoft, and Steve Weston. 2022. *Foreach: Provides Foreach Looping Construct*. <https://CRAN.R-project.org/package=foreach>.
- R Core Team. 2022. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Wickham, Hadley. 2014. "Tidy Data." *Journal of Statistical Software* 59 (10): 1–23. <https://doi.org/10.18637/jss.v059.i10>.
- Wickham, Hadley, Romain François, Lionel Henry, and Kirill Müller. 2022. *Dplyr: A Grammar of Data Manipulation*. <https://CRAN.R-project.org/package=dplyr>.